

```
In [1]: import warnings
warnings.filterwarnings("ignore", category=UserWarning)
import numpy as np
from tqdm import tqdm
import torch
import torch.nn as nn
import torchani
import matplotlib.pyplot as plt
```

```
In [2]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
#device = torch.device('cpu')
print(device)
```

cuda

```
In [3]: def init_aev_computer():
    Rcr = 5.2
    Rca = 3.5
    EtaR = torch.tensor([16], dtype=torch.float, device=device)
    ShfR = torch.tensor([
        0.900000, 1.168750, 1.437500, 1.706250,
        1.975000, 2.243750, 2.512500, 2.781250,
        3.050000, 3.318750, 3.587500, 3.856250,
        4.125000, 4.393750, 4.662500, 4.931250
    ], dtype=torch.float, device=device)

    EtaA = torch.tensor([8], dtype=torch.float, device=device)
    Zeta = torch.tensor([32], dtype=torch.float, device=device)
    ShfA = torch.tensor([0.90, 1.55, 2.20, 2.85], dtype=torch.float, device=device)
    ShfZ = torch.tensor([
        0.19634954, 0.58904862, 0.9817477, 1.37444680,
        1.76714590, 2.15984490, 2.5525440, 2.94524300
    ], dtype=torch.float, device=device)

    num_species = 4
    aev_computer = torchani.AEVComputer(
        Rcr, Rca, EtaR, ShfR, EtaA, Zeta, ShfA, ShfZ, num_species
    )
    return aev_computer

aev_computer = init_aev_computer()
aev_dim = aev_computer.aev_length
print(aev_dim)
```

384

```
In [4]: class AtomicNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(384, 128),
            nn.ReLU(),
            nn.Linear(128, 1)
        )

    def forward(self, x):
        return self.layers(x)
```

```
net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()
```

```
In [5]: def load_ani_dataset(dspath):
        self_energies = torch.tensor([
            0.500607632585, -37.8302333826,
            -54.5680045287, -75.0362229210
        ], dtype=torch.float, device=device)
        energy_shifter = torchani.utils.EnergyShifter(None)
        species_order = ['H', 'C', 'N', 'O']

        dataset = torchani.data.load(dspath)
        dataset = dataset.subtract_self_energies(energy_shifter, species_order)
        dataset = dataset.species_to_indices(species_order)
        dataset = dataset.shuffle()
        return dataset

dataset = load_ani_dataset("./ani_gdb_s01_to_s04.h5")
# Use dataset.split method to do split
train_data, val_data, test_data = dataset.split(.8,.1,.1)
```

```
In [6]: class ANITrainer:
        def __init__(self, model, batch_size, learning_rate, epoch, l2):
            self.model = model

            num_params = sum(item.numel() for item in model.parameters())
            print(f"{model.__class__.__name__} - Number of parameters: {num_params}")

            self.batch_size = batch_size
            self.optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
            self.epoch = epoch

        def train(self, train_data, val_data, early_stop=True, draw_curve=True):
            ### Eric's comment: here you should pass in train_data, val_data, not

            self.model.train()

            # init data loader
            print("Initialize training data...")
            ### Eric's comment: call the collate().cache() here to init data loader
            train_data_loader = train_data_loader = train_data.collate(batch_size)

            # definition of loss function: MSE is a good choice!
            loss_func = torch.nn.MSELoss()

            # record epoch losses
            train_loss_list = []
            val_loss_list = []
            lowest_val_loss = np.inf

            for i in tqdm(range(self.epoch), leave=True):
                train_epoch_loss = 0.0
                for train_data_batch in train_data_loader:

                    #computer energies
```

```

species = train_data_batch['species'].to(device)
coords = train_data_batch['coordinates'].to(device)
true_energies = train_data_batch['energies'].to(device).float(
_, pred_energies = model((species, coords))

#compute loss
batch_loss = loss_func(true_energies, pred_energies)

# do a step
### Eric's comment: here you need to do optimization, follow the
self.optimizer.zero_grad()
batch_loss.backward()
self.optimizer.step()

batch_importance = len(train_data_batch) / len(train_data)

### Eric's comment: instead of directly using batch_loss, please
### batch_loss.detach().cpu().item(), please refer to the previous
train_epoch_loss += batch_loss.detach().cpu().item() * batch_importance

# use the self.evaluate to get loss on the validation set
val_epoch_loss = self.evaluate(val_data, draw_plot=False)

# append the losses
### Eric's comment: train_epoch_loss should not be divided by len(train_data)
### because it is already multiplied by the batch_importance
train_loss_list.append(train_epoch_loss)
val_loss_list.append(val_epoch_loss)

if early_stop:
    if val_epoch_loss < lowest_val_loss:
        lowest_val_loss = val_epoch_loss
        weights = self.model.state_dict()

if draw_curve:
    x_axis = np.arange(self.epoch)
    fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)
    ax.set_yscale("log")
    # Plot train loss and validation loss
    ax.plot(x_axis, train_loss_list, label='Train')
    ax.plot(x_axis, val_loss_list, label='Validation')
    ax.legend()
    ax.set_xlabel("# Epoch")
    ax.set_ylabel("Loss")

if early_stop:
    self.model.load_state_dict(weights)

return train_loss_list, val_loss_list

def evaluate(self, data, draw_plot=True):

    # init data loader
    ### Eric's comment: again, call the collate().cache() here to init data loader
    data_loader = data.collate(batch_size).cache()

    # init loss function

```

```

loss_func = torch.nn.MSELoss()
total_loss = 0.0

if draw_plot:
    true_energies_all = []
    pred_energies_all = []

with torch.no_grad():
    for batch_data in data_loader:
        ### Eric's comment: here the name train_data_batch is not appropriate
        ### necessarily not train data

        #compute energies
        species = batch_data['species'].to(device)
        coords = batch_data['coordinates'].to(device)
        true_energies = batch_data['energies'].to(device).float()
        _, pred_energies = model((species, coords))

        #computer loss
        batch_loss = loss_func(true_energies, pred_energies)

        ### Eric's comment: here should be len(data) because the argument
        ### is called data, not train_data
        batch_importance = len(batch_data) / len(data)
        ### Eric's comment: again, instead of directly using batch_loss
        ### batch_loss.detach().cpu().item(), please refer to the previous
        total_loss += batch_loss.detach().cpu().item() * batch_importance

    if draw_plot:
        true_energies_all.append(true_energies.detach().cpu().numpy())
        pred_energies_all.append(pred_energies.detach().cpu().numpy())

if draw_plot:
    true_energies_all = np.concatenate(true_energies_all)
    pred_energies_all = np.concatenate(pred_energies_all)
    # Report the mean absolute error
    # The unit of energies in the dataset is hartree
    # please convert it to kcal/mol when reporting the mean absolute error
    # 1 hartree = 627.5094738898777 kcal/mol
    # MAE = mean(|true - pred|)
    hartree2kcalmol = 627.5094738898777
    mae = np.mean(np.abs((true_energies_all - pred_energies_all) * hartree2kcalmol))
    fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)
    ax.scatter(true_energies_all, pred_energies_all, label=f"MAE: {mae}")
    ax.set_xlabel("Ground Truth")
    ax.set_ylabel("Predicted")
    xmin, xmax = ax.get_xlim()
    ymin, ymax = ax.get_ylim()
    vmin, vmax = min(xmin, ymin), max(xmax, ymax)
    ax.set_xlim(vmin, vmax)
    ax.set_ylim(vmin, vmax)
    ax.plot([vmin, vmax], [vmin, vmax], color='red')
    ax.legend()

return total_loss

```

```

In [7]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
        model = nn.Sequential(
            aev_computer,
            ani_net

```

```
).to(device)

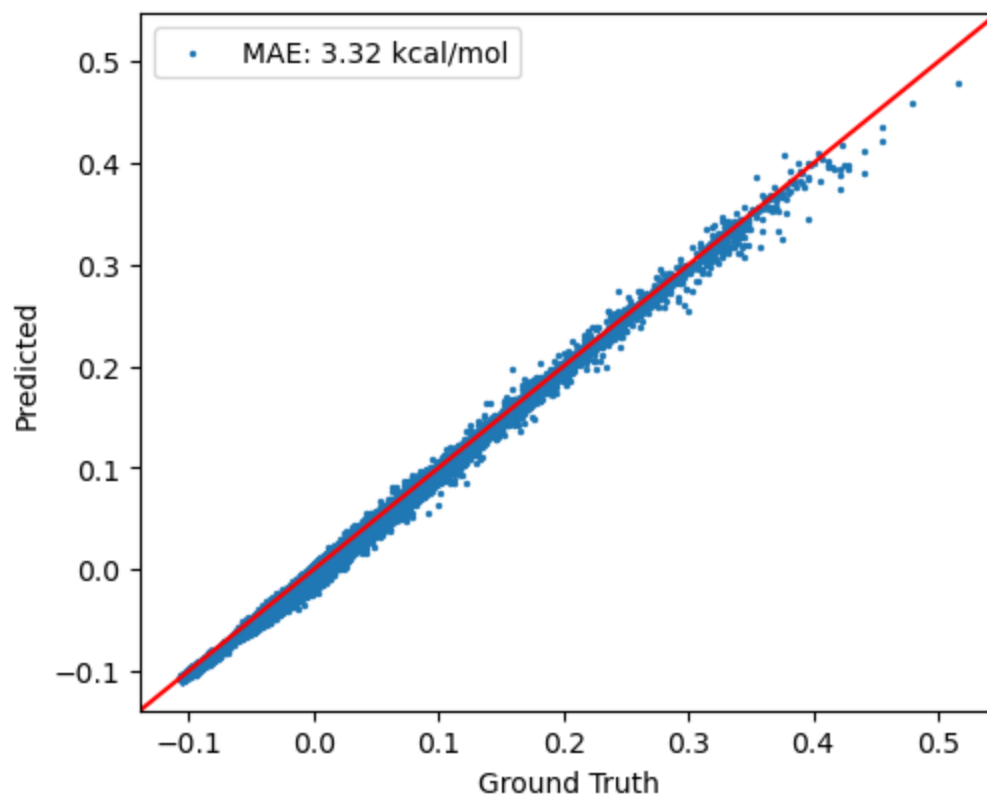
learning_rate = 1e-3
num_epochs = 30
l2 = 0.0
batch_size = 8162


trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 30/30 [05:11<00:00, 10.37s/it]





```
In [8]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)



learning_rate = 1e-4
num_epochs = 30
l2 = 0.0
batch_size = 8162

trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 30/30 [05:06<00:00, 10.23s/it]



```
In [9]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

learning_rate = 1e-2
num_epochs = 30
```


```
l2 = 0.0
batch_size = 8162

trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

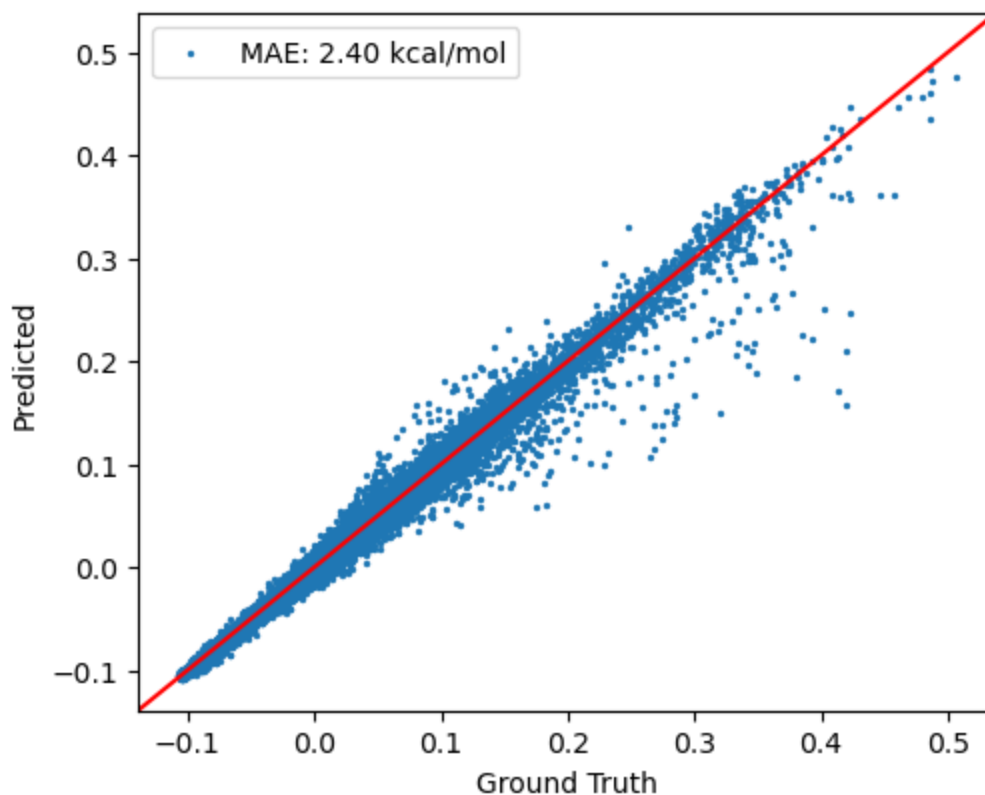
Sequential - Number of parameters: 197636

Initialize training data...

100%|██████████| 30/30 [04:55<00:00, 9.87s/it]







```
In [7]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)



learning_rate = 1e-1
num_epochs = 30
l2 = 0.0
batch_size = 8162

trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 30/30 [04:59<00:00, 10.00s/it]



```
In [9]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
        model = nn.Sequential(
            aev_computer,
            ani_net
        ).to(device)

        learning_rate = 1e-3
        num_epochs = 30
```


```
l2 = 1e-3
batch_size = 8162


trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100%|██████████| 30/30 [04:52<00:00, 9.76s/it]





```
In [10]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)



learning_rate = 1e-4
num_epochs = 30
l2 = 1e-3
batch_size = 8162

trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 30/30 [04:52<00:00, 9.76s/it]



```
In [11]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

learning_rate = 1e-4
num_epochs = 30
```


```
l2 = 0
batch_size = 8162


trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 30/30 [04:57<00:00, 9.91s/it]





```
In [12]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

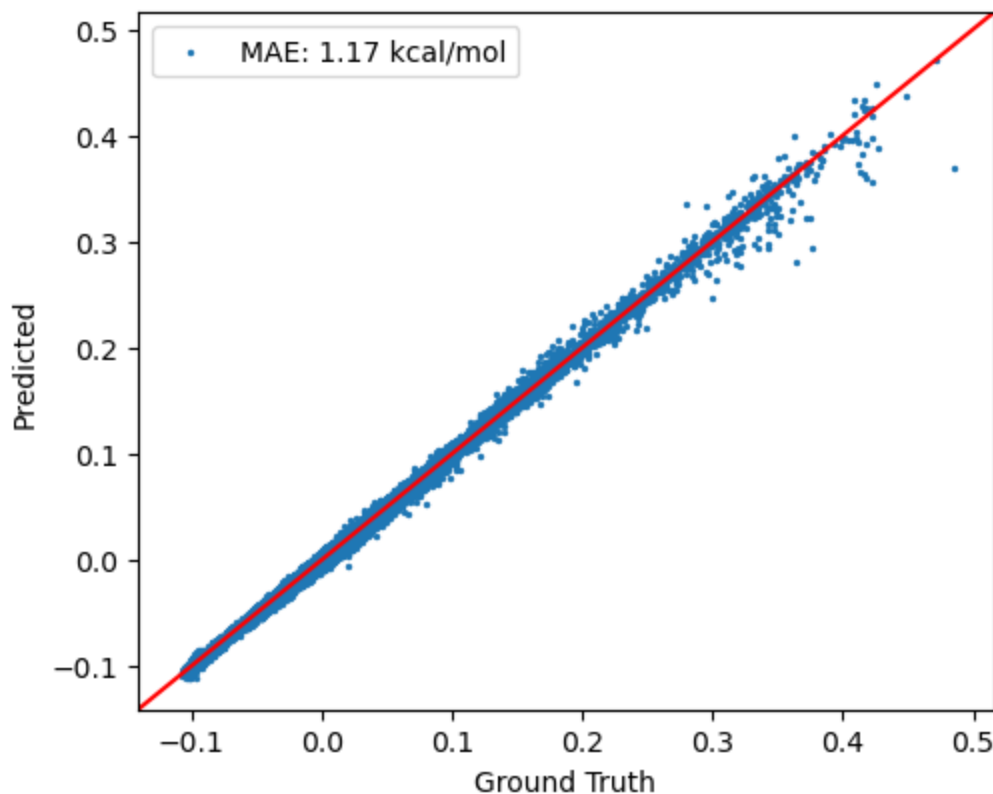

learning_rate = 1e-5
num_epochs = 50
l2 = 0
batch_size = 8162

trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 50/50 [07:57<00:00, 9.56s/it]



```
In [13]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

learning_rate = 1e-5
num_epochs = 50
```



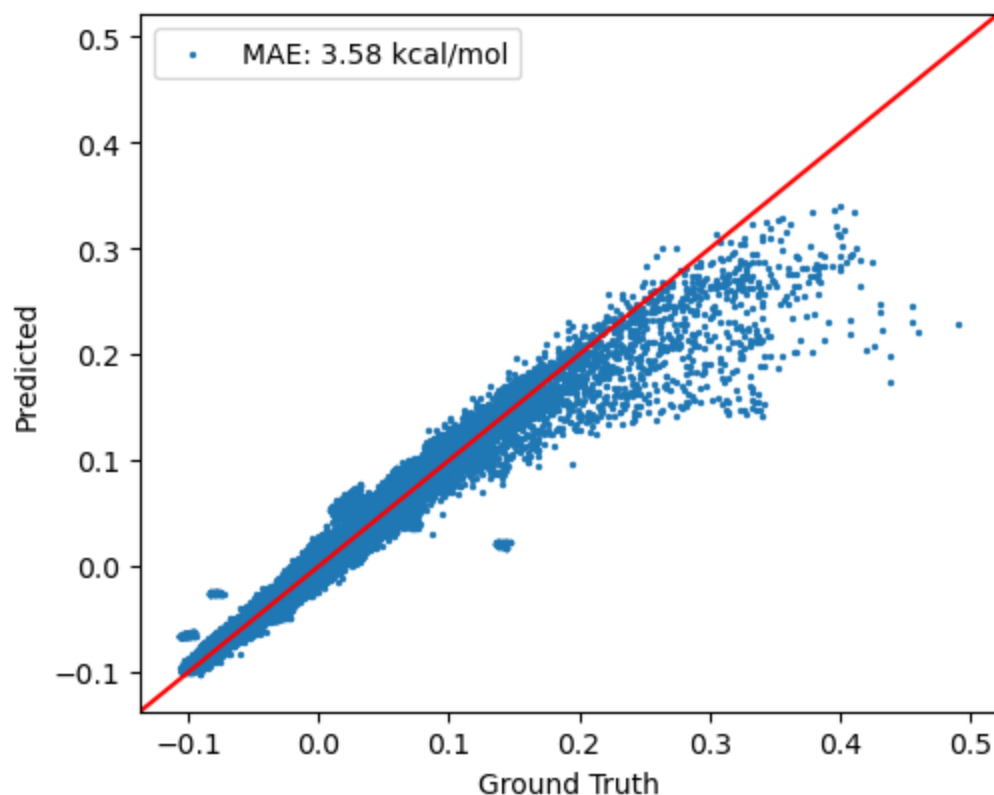
```
l2 = 1e-3
batch_size = 8162


trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 50/50 [08:16<00:00, 9.94s/it]





```
In [14]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)



learning_rate = 1e-5
num_epochs = 70
l2 = 0
batch_size = 8162

trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 70/70 [11:27<00:00, 9.81s/it]



```
In [15]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

learning_rate = 1e-5
num_epochs = 50
```


```
l2 = 1e-3
batch_size = 8162


trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 50/50 [08:02<00:00, 9.64s/it]





```
In [16]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)



learning_rate = 1e-4
num_epochs = 50
l2 = 0
batch_size = 8162

trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 50/50 [08:04<00:00, 9.69s/it]



```
In [17]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

learning_rate = 1e-4
num_epochs = 50
```


```
l2 = 1e-3
batch_size = 8162


trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100%|██████████| 50/50 [08:18<00:00, 9.98s/it]





```
In [10]: class AtomicNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(384, 256),
            nn.ReLU(),
            nn.Linear(256, 128),
            nn.ReLU(),
            nn.Linear(128, 1)
        )

    def forward(self, x):
        return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()

ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

learning_rate = 1e-3
num_epochs = 30
l2 = 0
batch_size = 8162

trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=False)
```





```
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```

Sequential – Number of parameters: 526340

Initialize training data...

100% |██████████| 30/30 [05:59<00:00, 11.99s/it]



```
In [11]: class AtomicNet(nn.Module):
         def __init__(self):
```

```

    super().__init__()
    self.layers = nn.Sequential(
        nn.Linear(384, 288),
        nn.ReLU(),
        nn.Linear(288, 192),
        nn.ReLU(),
        nn.Linear(192, 96),
        nn.ReLU(),
        nn.Linear(96, 1)
    )

    def forward(self, x):
        return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()

ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

learning_rate = 1e-3
num_epochs = 30
l2 = 0
batch_size = 8162



trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=False)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)

```

Sequential – Number of parameters: 739972

Initialize training data...

100%|██████████| 30/30 [06:45<00:00, 13.51s/it]



```
In [12]: class AtomicNet(nn.Module):
def __init__(self):
    super().__init__()
    self.layers = nn.Sequential(
        nn.Linear(384, 256),
        nn.ReLU(),
        nn.Linear(256, 128),
        nn.ReLU(),
```

```

        nn.Linear(128, 1)
    )

    def forward(self, x):
        return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()

ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

learning_rate = 1e-3
num_epochs = 30
l2 = 0
batch_size = 8162


trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=False)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)



```

Sequential – Number of parameters: 526340

Initialize training data...

100% |██████████| 30/30 [05:40<00:00, 11.36s/it]





```
In [15]: class AtomicNet(nn.Module):
def __init__(self):
    super().__init__()
    self.layers = nn.Sequential(
        nn.Linear(384, 128),
        nn.ReLU(),
        nn.Linear(128, 1)
    )
```


```
def forward(self, x):  
    return self.layers(x)
```



```
In [16]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)  
model = nn.Sequential(  
    aev_computer,  
    ani_net  
) .to(device)  
  
learning_rate = 1e-4  
num_epochs = 30  
l2 = 0.0  
batch_size = 8162  
  
trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)  
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)  
mae_val = trainer.evaluate(val_data)  
mae_test = trainer.evaluate(test_data)
```

Sequential – Number of parameters: 526340

Initialize training data...

100% |██████████| 30/30 [05:37<00:00, 11.26s/it]





In [ ]:

```
In [18]: class AtomicNet(nn.Module):  
         def __init__(self):  
             super().__init__()  
             self.layers = nn.Sequential(  
                 nn.Linear(384, 128),
```

```
        nn.ReLU(),
        nn.Linear(128, 1)
    )

    def forward(self, x):
        return self.layers(x)

ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)


learning_rate = 1e-3
num_epochs = 30
l2 = 0.0
batch_size = 8162

trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)
mae_val = trainer.evaluate(val_data)
mae_test = trainer.evaluate(test_data)
```



Sequential - Number of parameters: 526340

Initialize training data...

100% |██████████| 30/30 [05:50<00:00, 11.68s/it]







```
In [8]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)


learning_rate = 1e-4
num_epochs = 90
```



```
l2 = 1e-3  
batch_size = 8162  
  
trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)  
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)  
mae_val = trainer.evaluate(val_data)  
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 90/90 [14:44<00:00, 9.83s/it]





```
In [9]: class AtomicNet(nn.Module):
        def __init__(self):
            super().__init__()
            self.layers = nn.Sequential(
                nn.Linear(384, 128),
                nn.ReLU(),
                nn.Linear(128, 1)
            )
```

```
def forward(self, x):  
    return self.layers(x)  
  
net_H = AtomicNet()  
net_C = AtomicNet()  
net_N = AtomicNet()  
net_O = AtomicNet()  
  
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O]).to(device)  
model = nn.Sequential(  
    aev_computer,  
    ani_net  
).to(device)  
  
learning_rate = 1e-4  
num_epochs = 50  
l2 = 1e-4  
batch_size = 8162  
  
trainer = ANITrainer(model, batch_size, learning_rate, num_epochs, l2)  
train_losses, val_losses = trainer.train(train_data, val_data, early_stop=True)  
mae_val = trainer.evaluate(val_data)  
mae_test = trainer.evaluate(test_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100% |██████████| 50/50 [07:50<00:00, 9.42s/it]

